



POLITECNICO
MILANO 1863

**SCUOLA DI INGEGNERIA INDUSTRIALE
E DELL'INFORMAZIONE**

EXECUTIVE SUMMARY OF THE THESIS

The DRPM Strikes Back: More Flexibility for a Bayesian Spatio-Temporal Clustering Model

LAUREA MAGISTRALE IN MATHEMATICAL ENGINEERING - INGEGNERIA MATEMATICA

Author: FEDERICO ANGELO MOR

Advisor: PROF. ALESSANDRA GUGLIELMI

Co-advisor: ALESSANDRO CARMINATI

Academic year: 2023-2024

1. Introduction

Clustering is a fundamental technique of unsupervised learning where a set of data points has to be divided into homogenous groups of units which exhibit a similar behaviour relative to a target variable. It has long been a powerful tool for identifying structures and patterns within data, gaining increased popularity across a variety of scientific fields including social sciences, climate and environmental analysis, economics, and healthcare. Nowadays, as the accessibility of complex datasets has increased, the prevailing approach to tackling clustering challenges has transitioned from traditional algorithmic methods—such as hierarchical, partition-based, or density-based methods—to model-based methods, which assume a probabilistic modelling of the data. Thus, the development of these models finds its natural context within the Bayesian framework.

The effectiveness of Bayesian clustering models is particularly evident in the complex scenarios of spatio-temporal data, in which observations are collected over time and across various spatial locations. This complexity arises from the interaction between spatial and temporal dimensions; a complexity that is further compounded when covariates are also available. However, Bayesian mod-

els facilitate the incorporation of prior information into the model, thereby enhancing both flexibility of the formulation and interpretability of the results.

Over the years, several models have been proposed in the Bayesian literature. Among these, the Dependent Random Partition Model (Page et al., 2022) stands out for explicitly addressing the temporal dependence of partitions into the model formulation. However, the current implementation of its MCMC algorithm, written in C and accessible via an R interface, lacks some significant features, which comprise the inclusion of covariates, the handling of missing data, and a more efficient implementation.

In this work, we tackled these three shortcomings by enhancing the flexibility of the original model, yielding improved performances from both the statistical and computational perspectives.

2. Description of the model

The primary clustering tool employed in Bayesian clustering methods is a mixture model, which can be naturally extended to an infinite-mixture model within the framework of Bayesian nonparametrics. This extension specifies the conditional distribution of data points given a realization of the partition of the units, and a prior is assigned to

this partition by means of the discreteness of the assigned random probability measure (RPM). In the spatio-temporal context, where we analyse the sequence of partitions ρ_1, \dots, ρ_T , i.e. the clusters of the n experimental units at each time instant $t = 1, \dots, T$, this approach may not always be optimal. In fact, in infinite mixture models, partitions are only *induced* by the random partition model; however, when the sequence of partitions is the actual inferential object of interest, it is essential to model this sequence *directly* to accurately capture the temporal dependencies inherent in the partitions. Although many Bayesian nonparametric methods have been developed to temporally correlate a sequence of random probability measures, they typically implement this temporal dependence through the atoms or weights associated with the representation of the discrete RPM. However, the correlations modelled in these underlying parameters does not necessarily reflect into correlations among the induced partitions. In contrast, the Dependent Random Partition Model (DRPM) (Page et al., 2022) explicitly models temporal dependencies among the partitions within its formulation. This feature enhances the identification of temporal trends, resulting in clusters that evolve in a more gentle and interpretable way.

2.1. Temporal modelling for sequences of partitions

We now describe how the DRPM models temporal relationships within sequences of partitions. Before proceeding with the explanation, we first establish some general notation. We consider a spatio-temporal context where there are n units to be clustered at all time points $t = 1, \dots, T$. Each unit is represented as $i = 1, \dots, n$. We denote by $\rho_t = \{S_{1t}, \dots, S_{k_t t}\}$ the partition at time t of the n experimental units, composed by k_t clusters. An alternative representation of this partition is possible through cluster labels $\mathbf{c}_t = \{c_{1t}, \dots, c_{nt}\}$, with $c_{it} = j$ indicating that unit i belongs to cluster S_{jt} . Finally, we will denote with a \star superscript all the variables or quantities which are cluster-specific. Introducing temporal dependence in a sequence of partitions requires the formulation of a joint probability model for (ρ_1, \dots, ρ_T) . Temporal dependence among the ρ_t 's implies that the cluster configuration in ρ_t may be influenced by the cluster configurations found in $\rho_{t-1}, \rho_{t-2}, \dots, \rho_1$. However, Page et al. (2022) restricted this temporal connection to a first-order Markovian struc-

ture, where the conditional distribution of ρ_t given all the predecessors $\rho_{t-1}, \rho_{t-2}, \dots, \rho_1$ actually depends only on ρ_{t-1} . This led to the construction of the joint probability model for (ρ_1, \dots, ρ_T) as

$$P(\rho_1, \dots, \rho_T) = P(\rho_T | \rho_{T-1}) \cdots P(\rho_2 | \rho_1) P(\rho_1) \quad (1)$$

Here, $P(\rho_1)$ is an exchangeable partition probability function (EPPF) which describes how the n experimental units at time period 1 are grouped into k_1 distinct clusters. For this function, Page et al. (2022) employed a Chinese restaurant process, in the form

$$P(\rho_1 | M) \propto \prod_{j=1}^{k_1} M \cdot (|S_j| - 1)! \quad (2)$$

In Section 2.2 we will detail how this EPPF can be adapted to incorporate spatial and covariates information.

To characterize the other terms $P(\rho_t | \rho_{t-1})$ in (1), the following auxiliary variables need to be introduced for all units $i = 1, \dots, n$.

$$\gamma_{it} = \begin{cases} 1 & \text{if unit } i \text{ is not reallocated when} \\ & \text{moving from time } t-1 \text{ to } t \\ 0 & \text{otherwise} \end{cases}$$

These parameters model the similarity between ρ_{t-1} and ρ_t . If the partitions ρ_{t-1} and ρ_t are highly dependent, their cluster configurations will change minimally, resulting in only few units changing their cluster assignments. Conversely, partitions exhibiting low dependence will likely manifest significantly different cluster configurations, leading to a majority of the units being reallocated. By construction, we set $\gamma_{i1} = 0$ for all i , meaning that at the first time instant all units will get reallocated as there is no partition at time $t = 0$. Page et al. (2022) assumed $\gamma_{it} \stackrel{\text{ind}}{\sim} \text{Ber}(\alpha_t)$ where $\alpha_t \in [0, 1]$ serves as a temporal dependence parameter. Specifically, $\alpha_t = 1$ denotes perfect temporal dependence, with $\rho_t = \rho_{t-1}$, while $\alpha_t = 0$ implies full independence of ρ_t from ρ_{t-1} . For clarity, the vector $\boldsymbol{\gamma}_t = (\gamma_{1t}, \dots, \gamma_{nt})$ is introduced, so that the T pairs of parameters $(\boldsymbol{\gamma}_j, \rho_j)$ are explicitly reported in the augmented formulation of the joint model (1), which becomes

$$P(\boldsymbol{\gamma}_1, \rho_1, \dots, \boldsymbol{\gamma}_T, \rho_T) = P(\rho_T | \boldsymbol{\gamma}_T, \rho_{T-1}) P(\boldsymbol{\gamma}_T) \cdots P(\rho_2 | \boldsymbol{\gamma}_2, \rho_1) P(\boldsymbol{\gamma}_2) P(\rho_1) \quad (3)$$

Once the partition model is specified, there is considerable flexibility in how to define the remainder

of the Bayesian model. The complete formulation designed by Page et al. (2022) is provided in (4),

$$\begin{aligned}
Y_{it}|Y_{it-1}, \boldsymbol{\mu}_t^*, \boldsymbol{\sigma}_t^{2*}, \boldsymbol{\eta}, \mathbf{c}_t &\stackrel{\text{ind}}{\sim} \mathcal{N}(\mu_{c_{it}}^* + \eta_{1i}Y_{it-1}, \sigma_{c_{it}t}^{2*}(1 - \eta_{1i}^2)) \\
Y_{i1} &\stackrel{\text{ind}}{\sim} \mathcal{N}(\mu_{c_{i1}}^*, \sigma_{c_{i1}}^{2*}) \\
\xi_i = \text{Logit}(\tfrac{1}{2}(\eta_{1i} + 1)) &\stackrel{\text{ind}}{\sim} \text{Laplace}(a, b) \\
(\mu_{jt}^*, \sigma_{jt}^{2*}) &\stackrel{\text{ind}}{\sim} \mathcal{N}(\vartheta_t, \tau_t^2) \times \mathcal{U}(0, A_\sigma) \\
\vartheta_t|\vartheta_{t-1} &\stackrel{\text{ind}}{\sim} \mathcal{N}((1 - \varphi_1)\varphi_0 + \varphi_1\vartheta_{t-1}, \lambda^2(1 - \varphi_1^2)) \\
(\vartheta_1, \tau_1) &\stackrel{\text{iid}}{\sim} \mathcal{N}(\varphi_0, \lambda^2) \times \mathcal{U}(0, A_\tau) \\
(\varphi_0, \varphi_1, \lambda) &\sim \mathcal{N}(m_0, s_0^2) \times \mathcal{U}(-1, 1) \times \mathcal{U}(0, A_\lambda) \\
\{\mathbf{c}_t, \dots, \mathbf{c}_T\} &\sim \text{tRPM}(\boldsymbol{\alpha}, M) \text{ with } \alpha_t \stackrel{\text{iid}}{\sim} \text{Beta}(a_\alpha, b_\alpha)
\end{aligned} \tag{4}$$

where Y_{it} denotes the response variable measured by the i -th unit at time t and $\text{tRPM}(\boldsymbol{\alpha}, M)$ represents the temporal random partition model (3) parametrised by $\alpha_1, \dots, \alpha_T$ and the EPPF in (2). Our generalization of the original formulation is presented in (5), with changes and additions that differ from the original model (4) highlighted in light blue.

$$\begin{aligned}
Y_{it}|Y_{it-1}, \boldsymbol{\mu}_t^*, \boldsymbol{\sigma}_t^{2*}, \boldsymbol{\eta}, \mathbf{c}_t &\stackrel{\text{ind}}{\sim} \mathcal{N}(\mu_{c_{it}}^* + \eta_{1i}Y_{it-1} + \mathbf{x}_{it}^T \boldsymbol{\beta}_t, \sigma_{c_{it}t}^{2*}(1 - \eta_{1i}^2)) \\
Y_{i1} &\stackrel{\text{ind}}{\sim} \mathcal{N}(\mu_{c_{i1}}^* + \mathbf{x}_{i1}^T \boldsymbol{\beta}_1, \sigma_{c_{i1}}^{2*}) \\
\boldsymbol{\beta}_t &\stackrel{\text{ind}}{\sim} \mathcal{N}_p(\mathbf{b}, s^2 \mathbf{I}) \\
\xi_i = \text{Logit}(\tfrac{1}{2}(\eta_{1i} + 1)) &\stackrel{\text{ind}}{\sim} \text{Laplace}(a, b) \\
(\mu_{jt}^*, \sigma_{jt}^{2*}) &\stackrel{\text{ind}}{\sim} \mathcal{N}(\vartheta_t, \tau_t^2) \times \text{invGamma}(a_\sigma, b_\sigma) \\
\vartheta_t|\vartheta_{t-1} &\stackrel{\text{ind}}{\sim} \mathcal{N}((1 - \varphi_1)\varphi_0 + \varphi_1\vartheta_{t-1}, \lambda^2(1 - \varphi_1^2)) \\
(\vartheta_1, \tau_1^2) &\stackrel{\text{iid}}{\sim} \mathcal{N}(\varphi_0, \lambda^2) \times \text{invGamma}(a_\tau, b_\tau) \\
(\varphi_0, \varphi_1, \lambda^2) &\sim \mathcal{N}(m_0, s_0^2) \times \mathcal{U}(-1, 1) \times \text{invGamma}(a_\lambda, b_\lambda) \\
\{\mathbf{c}_t, \dots, \mathbf{c}_T\} &\sim \text{tRPM}(\boldsymbol{\alpha}, M) \text{ with } \alpha_t \stackrel{\text{iid}}{\sim} \text{Beta}(a_\alpha, b_\alpha)
\end{aligned} \tag{5}$$

In our updated formulation, we opted to model the variance parameters using an inverse gamma distribution instead of the the uniform distribution employed in (4). This more sophisticated choice should ensure better mixing in the Markov chain, as the inverse gamma distributions recover conjugacy within the model, allowing for variance updates to be performed using the analytically exact Gibbs sampler rather than relying on the acceptance-rejection method of Metropolis algorithm. Moreover, we introduced the regression parameter $\boldsymbol{\beta}_t$, which incorporates covariates in the likelihood, with the goal of improving the accuracy in the fitted estimates of the target variable.

2.2. MCMC algorithm

We now detail the MCMC algorithm developed by Page et al. (2022), which is necessary for sampling from the posterior distributions implied by model (4). The MCMC algorithm associated to our generalized model (5) required only minor ad-

justments to the original structure.

For the sake of clarity, we will refer to CDRPM for the original model formulation by (Page et al., 2022), and to JDRPM for our updated version. The letters C and J denote the programming languages used to implement their corresponding MCMC algorithms: C for the former, Julia for the latter.

The iterative structure of (3) suggests the use of a Gibbs sampler, where γ_t and ρ_t are updated sequentially. The Markovian assumption reduces computational costs as we only need to consider ρ_{t-1} and ρ_{t+1} when updating ρ_t . To perform this update, the terms $P(\rho_1)$ and $P(\rho_t|\rho_{t-1})$ of (3) need to be clarified. While $P(\rho_1)$ is defined by (2), the derivation of $P(\rho_t|\rho_{t-1})$ requires the concept of compatibility.

Definition 2.1 (compatibility). Two partitions ρ_t and ρ_{t-1} are said to be *compatible* with respect to γ_t if ρ_t can be obtained from ρ_{t-1} by reallocating items as indicated by γ_t ; that is, by moving the units i for which $\gamma_{it} = 0$.

To perform this compatibility check, it suffices to ensure that the reduced partitions from ρ_t and ρ_{t-1} are identical, where “reduced” refers to the restriction of partitions ρ_t and ρ_{t-1} to the units that cannot move. Indeed, if these fixed units are clustered in the same way, then the free movers from ρ_t can be assigned labels to match the ones assigned in ρ_{t-1} . Denoting the set of fixed units at time t as $\mathfrak{R}_t = \{i : \gamma_{it} = 1\}$, this check translates into verifying that $\rho_t^{\mathfrak{R}_t} = \rho_{t-1}^{\mathfrak{R}_t}$.

This compatibility must be verified during the update steps of parameters γ_{it} and cluster labels c_{it} to ensure that the new sampled draws remain valid and coherent across all partitions and parameters involved. To derive $P(\rho_t|\rho_{t-1})$, Page et al. (2022) proved that, under the construction outlined thus far, marginally ρ_1, \dots, ρ_T are identically distributed with law derived from the EPPF used to model ρ_1 . This result, along with the previously outlined compatibility analysis, enables the derivation of the complete update rules for γ_{it} and c_{it} .

Once the sampling scheme for γ_{it} and ρ_t is established, updating the other parameters of the model becomes straightforward. Each parameter will either require a Gibbs step or a Metropolis step, depending on the conjugacy of the prior associated with that parameter. Thus, we will now outline the final theoretical step, which addresses

how spatial and covariates information can be included in the prior for the partitions.

The core of the clustering process, that is, updating γ_{it} and ρ_t , is inherently complex as it necessitates checking for compatibility issues. The approach entails simulating the assignment of each unit i , currently belonging to cluster j , to either one of the existing clusters or to a new singleton cluster. For each scenario, we compute the probability of this assignment to occur, from which we derive weights to inform the sampling decision for the next iteration. To influence the definition of these weights by spatial and, in our JDRPM updated formulation, covariates information, the EPPF (2) has to be updated through dedicated cohesion and similarity functions, respectively.

2.3. Spatial cohesions

The incorporation of spatial information can be effectively accommodated through the EPPF in our framework, resulting in spatially informed clusters that evolve over time. To introduce this extension, let \mathbf{s}_i denote the spatial coordinates of the i -th (we note that these coordinates do not change over time), and let \mathbf{s}_{jt}^* denote the subset of spatial coordinates of the units belonging to cluster S_{jt} . Then, we can express the EPPF for the t -th partition in the following product form

$$P(\rho_t|M, \mathcal{S}) \propto \prod_{j=1}^{k_t} C(S_{jt}, \mathbf{s}_{jt}^*|M, \mathcal{S}) \quad (6)$$

Compared to the original formulation of (2), where $P(\rho_t|M) \propto \prod_{j=1}^{k_t} c(S_{jt}|M)$, (6) incorporates a spatial component into the partition weights through the spatial cohesion function $C(S_{jt}, \mathbf{s}_{jt}^*|M, \mathcal{S})$. The original term $c(S_{jt}|M)$ describes how units inside cluster S_{jt} are likely to be clustered together a priori, while the cohesion function $C(S_{jt}, \mathbf{s}_{jt}^*|M, \mathcal{S})$, parametrised by a set of parameters \mathcal{S} , measures the compactness of the spatial coordinates \mathbf{s}_{jt}^* . In the JDRPM code, we implemented the same six cohesion functions which were also implemented in CDRPM, specifically, those described in (Page et al., 2016).

2.4. Covariates similarities

The incorporation of covariates information, a characteristic feature of our generalized model, can be similarly integrated into the EPPF definition. To introduce this extension, let X_{jt}^* denote the $p \times |S_{jt}|$ matrix that contains the p covariates of the units belonging to cluster S_{jt} , i.e.

$X_{jt}^* = \{\mathbf{x}_{it}^* = (x_{it1}, \dots, x_{itp})^T : i \in S_{jt}\}$. In the current implementation of JDRPM we decided to treat each covariate individually. In this way, the new term in the definition of the EPPF for $P(\rho_t)$ will be a function of the vector \mathbf{x}_{jtr}^* that collects the values of the r -th covariate for the units inside cluster S_{jt} , i.e. row r of matrix X_{jt}^* . Then, each contribution of the covariates will be considered independently, leading to an EPPF in the form

$$P(\rho_t|M, \mathcal{S}, \mathcal{C}) \propto \prod_{j=1}^{k_t} C(S_{jt}, \mathbf{s}_{jt}^*|M, \mathcal{S}) \left(\prod_{r=1}^p g(S_{jt}, \mathbf{x}_{jtr}^*|\mathcal{C}) \right) \quad (7)$$

Here, $g(S_{jt}, \mathbf{x}_{jtr}^*|\mathcal{C})$ is the covariate similarity, parametrised by a set of parameters \mathcal{C} , which measures the similarity of the covariates values \mathbf{x}_{jtr}^* . This approach of treating each covariate individually is convenient as it simplifies the definition of the similarity functions and allows to accommodate both numerical and categorical covariates. In the JDRPM code, we implemented four similarity functions, which are detailed in (Page et al., 2018).

3. Implementation

The MCMC algorithm to compute the posterior samples of our generalized model, described in Section 2.2, has been implemented in Julia (Bezanson et al., 2017).

Julia is a relatively new programming language that combines the ease and expressiveness of high-level languages with the efficiency and performance of low-level languages. This balance is primarily achieved through just-in-time compilation, using the LLVM framework, along with features as dynamic multiple dispatch and extensive code specialization against multiple run-time types.

This positions Julia as a highly effective language for improving productivity and performance in implementations. For reference, our development process comported an improvement by 74% in execution time and by 88% in the memory allocated, when comparing an early-stage version of the Julia code to the final version. This significant performance gain was mainly attributable to a more effective memory management.

In fact, during the earlier development phases, we observed that a substantial portion of execution time was not spent by actual calculations, but rather by Julia's garbage collector, which had the burden of tracking all the allocated memory and reclaiming the unused one to make it available for new computations. To address this issue,

we refactored the code to operate more in-place. This approach involved passing as arguments the variables that would be modified by a function, and applying those changes directly within the function, rather than returning values and subsequently using them to update the original variables. This optimization strategy was facilitated by Julia profiling tools such as `ProfileCanvas` and `BenchmarkTools`.

Another relevant performance improvement was achieved through the optimized implementation of spatial cohesions and covariates similarities. In fact, these functions are essential for deriving the partition weights in the update steps of γ_{it} and ρ_t , and as such they are potentially executed millions of times during each fit. Therefore, refining these functions was crucial for providing fast execution times. In the cohesions functions, optimal performance was reached by using static vectors and matrices, provided by the `StaticArrays` package, rather than standard dynamic structures. For the similarity functions, we employed a SIMD paradigm (Single Instruction Multiple Data), through the `@simd` macro, which accelerated the elaboration of the covariates values by working on multiple data-chunks at once, rather than processing each element individually.

4. Analysis of the models

To evaluate the numerical performance of both algorithms, we analysed posterior samples and clusters estimates in two scenarios: first using a synthetic dataset that includes only the response variable, and secondly employing a real-world spatio-temporal dataset. For the latter, we used a dataset from the `AgrImOnIA` project (Fassò et al., 2023) which encompasses weekly measurements of air pollutants in 2018, together with many other environmental variables, in the Lombardy region of Italy. This dataset also provides covariates, whose effects will be examined in Sections 4.1.1 and 4.1.2.

4.1. Statistical improvements

In the first set of experiments we fitted the two models, in both simulated and real-world scenarios, under equivalent conditions—i.e. identical datasets, hyperparameters, and MCMC configurations. With this setup, the models are expected to perform similarly to each other, as JDRPM is a generalization of CDRPM and should therefore manifest the same core behaviour of the orig-

inal model. Indeed, the clusters estimates, as well as their temporal trend depicted in Figures 1 and 2, were remarkably similar, thus confirming the correctness of the JDRPM implementation. Moreover, the fitted values were more accurate in the JDRPM fits, as demonstrated by the mean squared errors (MSEs) improved by 20% and 8% in the simulated and real-world scenarios, respectively, compared to CDRPM fits. In absence of additional information from covariates, this improvement was likely attributable to the better mixing properties of the Markov chain in our model, thanks to the distribution choice for the variance parameters.




Figure 1: Lagged ARI values of CDRPM and JDRPM fits, in the simulated data scenario.




Figure 2: Lagged ARI values of CDRPM and JDRPM fits, in the real-world scenario.

Next, we investigated the effects of covariates-informed fits. Given their different roles, we analysed separately covariates in the likelihood and covariates in the prior.

4.1.1 Covariates in the likelihood

The inclusion of covariates in the likelihood through the regression term β_t had the main goal of improving the accuracy of the fitted values. Therefore, we examined two natural use-cases of covariates in the likelihood: first, fitting on spatio-temporal dataset with missing data, and second, performing inference on a new unit at a new location, as in a typical kriging scenario.

In the first experiment, we performed JDRPM fits on an incomplete dataset, generated by randomly removing 10% of the data points from the full dataset. Indeed, the inclusion of covariates in the likelihood from the spatio-temporal dataset resulted in improvements across all metrics, as proved by the reduced MSEs and better fitting metrics LPML and WAIC. The results are summarized in Table 1.

Table 1: Comparison of JDRPM fits, in the real-world scenario, with and without the inclusion of covariates in the likelihood, on a complete dataset and on a dataset with missing values.

	Full data	Full data + Xlk	NA data	NA data + Xlk
MSE (mean)	0.0131	0.0112	0.0160	0.0127
MSE (median)	0.0138	0.0113	0.0170	0.0130
LPML	624.91	778.96	502.86	625.81
WAIC	-1898.05	-2029.84	-1793.64	-1902.74
Execution time	48m	56m	43m	58m

We also mention that the standard fits without covariates in the likelihood did not exhibit a significant drop in performance, compared to their full dataset counterparts. This indicates that even in the absence of additional information, the update rule for the missing data Y_{it} remains effective. For example, in the simulated scenario with missing data, where covariates information was not available, all the true values of the missing points fell within the 95% credible intervals of the corresponding fitted estimates.

Regarding the kriging analysis, we simulated the insertion of new units at new locations by removing all data entries from three randomly selected units within the spatio-temporal dataset. In this experiment, we aimed to evaluate the JDRPM’s accuracy in predicting the behaviour of units for which sensors may be absent or inactive, and we expected that the estimation accuracy would improve as model complexity increases. Indeed, the best performances, in terms of MSE, LPML, and WAIC, was achieved by JDRPM fits with covariates in the likelihood; proving again the effectiveness of this addition to the model formulation. However, some inaccuracies were observed with particular unit which was often clustered as a singleton, exhibited extreme values in some of the included covariates, and was quite separate from other units. Thus, these particular factors possibly hindered a more precise estimation of her response variable values.

4.1.2 Covariates in the prior

The main goal of incorporating covariates in the prior was to generate more informative and accurate clusters configurations. Therefore, in this set of experiments, we performed several JDRPM fits on the PM_{10} target variable, but now including three relevant covariates in the prior: wind speed, precipitation levels, and boundary layer height. Again, the inclusion of covariates information displayed a significant improvement across all metrics, as reported in Table 2.

Table 2: Comparison between CDRPM and JDRPM fits and their associated algorithms, in the real-world scenario, with and without covariates in the prior.

	CDRPM	JDRPM	JDRPM + Xcl
MSE (mean)	0.0142	0.0131	0.0126
MSE (median)	0.0149	0.0138	0.0135
LPML	694.81	624.91	677.71
WAIC	-1768.42	-1898.05	-1969.76
Execution time	1h38m	48m	1h20m

Although incorporating multiple covariates may lead to divergent “clustering suggestions,” their contributions can still be effectively assessed through cluster-specific boxplots of the target variable PM_{10} alongside the three covariates included in the analysis. Figure 3 demonstrates how the spatially and covariates-informed JDRPM fit captured a more refined and coherent clustering structure, providing a clear separation in the target variable PM_{10} , as well as among the covariates, particularly regarding wind speed and total precipitation. Such separation facilitates a clearer interpretation of the clustering results. For instance, the dark red cluster is characterized by the highest levels of PM_{10} and the second-last levels of wind speed and total precipitations; factors which likely hindered the pollutant dispersion.

The contribution of covariates is particularly evident when comparing the CDRPM spatially-informed fit to the JDRPM fit including space and covariates in the prior. Such comparison, exemplified by Figures 4 and 5, demonstrates how the JDRPM fit was more effective in identifying clusters that appeared, instead, sparser and less specialized in the CDRPM fit.

4.2. Computational improvements

Finally, to rigorously assess whether the objective of faster execution times was achieved, we designed a series of experiments to compare the two

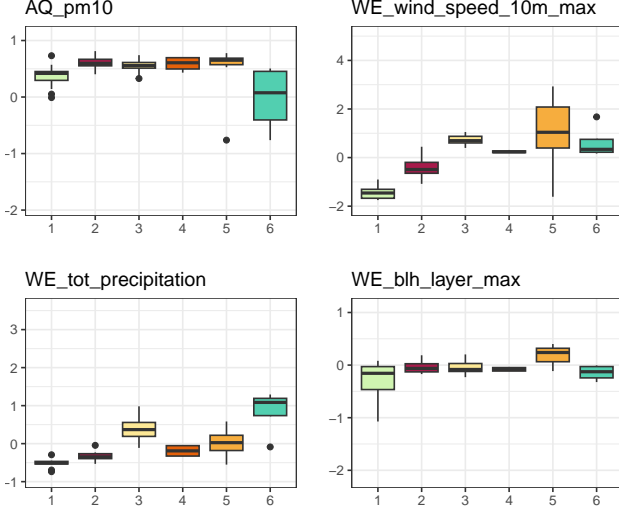


Figure 3: Example of the clusters estimates produced by JDRPM fit, in the real-world scenario, with covariates in the prior. The boxplots refer to the target variable AQ_pm10 along with the three covariates included in the fit.




Figure 4: Example of clusters estimates with respect to the wind speed covariate, in the CDRPM spatially-informed fit.




Figure 5: Example of clusters estimates with respect to the wind speed covariate, in the JDRPM spatially-informed fit with covariates in the prior.

models and their corresponding implementations. The varying levels of information influence model complexity, which in turn affects computational load; however, also the size of the testing dataset also plays a significant role. To systematically investigate these factors, we conducted experiments across incremental levels of information—starting with only the target variable, then adding spatial information, covariates in the likelihood, and finally covariates in the prior—as well as across a

“mesh” of dataset sizes, with both the number of units n and the time horizons T ranging through the set $\{10, 50, 100, 250\}$. It is important to note that results from datasets with extreme sizes, such as those with n or T being 250, may reflect hardware limitations of the machine used for fitting, i.e. my simple laptop, rather than actual model performance. Therefore, more attention should be devoted to the more accurate intermediate-sized experiments, which firmly demonstrate the JDRPM’s improved performance. For instance, Figure 6 proves that, in the case of spatially-informed fits, JDRPM can operate up to twice as fast as CDRPM.

fit J – target + space					ms/it
	n=10	n=50	n=100	n=250	
T=250	8.3 (1.64x)	96 (1.28x)	380 (1.01x)	2060 (0.98x)	2000
T=100	3.16 (1.68x)	28.8 (1.68x)	98.8 (1.54x)	563 (1.42x)	1500
T=50	1.5 (1.77x)	12.4 (1.94x)	41.6 (1.83x)	222.5 (1.81x)	1000
T=10	0.3 (1.79x)	2.28 (2x)	7.6 (1.95x)	42.2 (1.82x)	500
					0

Figure 6: Execution times, in milliseconds per iteration, when fitting JDRPM on a spatio-temporal dataset. In brackets are reported the speedups relative to the CDRPM timings, where higher values indicate better performance.

Moreover, JDRPM fits including covariates information still exhibit strong performance, despite the additional complexity and subsequent increase in computational load. Figure 7 illustrates that, on a medium sized dataset consisting of $n = 50$ units, fully-informed JDRPM fits can perform faster than spatially-informed CDRPM fits. In addition, we can observe how JDRPM fits with spatial information can perform faster than CDRPM fits without spatial information.

In Figure 7, the number of covariates in the likelihood and prior levels was fixed to $p = 5$. This configuration appears to represent the equilibrium point between JDRPM and CDRPM performances. More precisely, other experiments were conducted with varying numbers of covariates in likelihood and prior levels, demonstrating how in

the same amount of time that CDRPM performs a spatially-informed fit, JDRPM can perform a spatially-informed fit with up to five covariates in the prior and with an almost indifferent number of covariates in the likelihood, as covariates in the likelihood do not add significant computational load. These results strongly support the performance speedup achieved with our new implementation.




Figure 7: Visual representation of the performance of all the fits, for the $n = 50$ case. Specifically, the execution time per iteration of JDRPM spatially-informed fit has been taken as a reference, to which then all other fits have been compared to derive their speedup (or slowdown) factor. Points above the reference line indicate slower fits, while points below denote faster fits.

5. Conclusions

In conclusion, we can state that JDRPM represents a significant improvement over the original formulation of CDRPM. From a theoretical perspective, JDRPM retains the foundational structure of its predecessor, while also introducing covariates in the prior and likelihood levels, as well as providing more flexibility through the handling of missing data in the target variable. From a computational perspective, we successfully reduced execution times compared to the original implementation of Page et al. (2022). Moreover, the innovative choice of the Julia language is expected to facilitate the implementation of future modifications or updates.

However, potential drawbacks arise from the increased complexity of our generalized model. The choice of a more sophisticated distribution for the variance parameters, along with the selection of parameters for spatial cohesions and covariates similarities, may necessitate more careful consideration and experimental tuning when performing

a model fit. To this end, future development paths could focus on mitigating this complexity, e.g. through heuristics for assigning the various parameters and hyperparameters given the dataset at hand. Additionally, the Julia powerful ecosystem could be leveraged to provide real-time feedbacks on the sampling process, e.g. through live trace plots of the sampled values or logs detailing the progression of MCMC update steps.

In summary, while the JDRPM’s complexity may present certain challenges, it also establishes a foundation for significant methodological advancements and practical enhancements. Our developments are expected to improve the model’s applicability and accuracy, ultimately leading to more effective research outcomes.

References

- Bezanson, Jeff et al. (2017). “Julia: A fresh approach to numerical computing”. In: *SIAM Review* 59.1, pp. 65–98. DOI: 10.1137/141000671 (cit. on p. 4).
- Fassò, A. et al. (2023). *AgrImOnIA: Open Access dataset correlating livestock and air quality in the Lombardy region, Italy (3.0.0)*. DOI: <https://doi.org/10.5281/zenodo.7956006> (cit. on p. 5).
- Page, Garritt L. and Fernando A. Quintana (Sept. 2018). “Calibrating covariate informed product partition models”. In: *Statistics and Computing* 28, pp. 1–23. DOI: 10.1007/s11222-017-9777-z (cit. on p. 4).
- (2016). “Spatial Product Partition Models”. In: *Bayesian Analysis* 11.1, pp. 265–298. DOI: 10.1214/15-BA971 (cit. on p. 4).
- Page, Garritt L., Fernando A. Quintana, and David B. Dahl (2022). “Dependent Modeling of Temporal Sequences of Random Partitions”. In: *Journal of Computational and Graphical Statistics* 31.2, pp. 614–627. DOI: 10.1080/10618600.2021.1987255 (cit. on pp. 1–3, 8).